

CPU cache

Acknowledgment

These slides are based on the slides of Prof. Weatherspoon of Cornell University and Dr Manoharan of Auckland University

1

Recommended Reading

Patterson & Hennessy, Computer Organization & Design: The Hardware/Software Interface, Morgan Kaufmann, 5th edition (Chapter 5.1-5.3 and any other sections that appear relevant)
(4th edition, 3rd or 2nd edition are also OK)

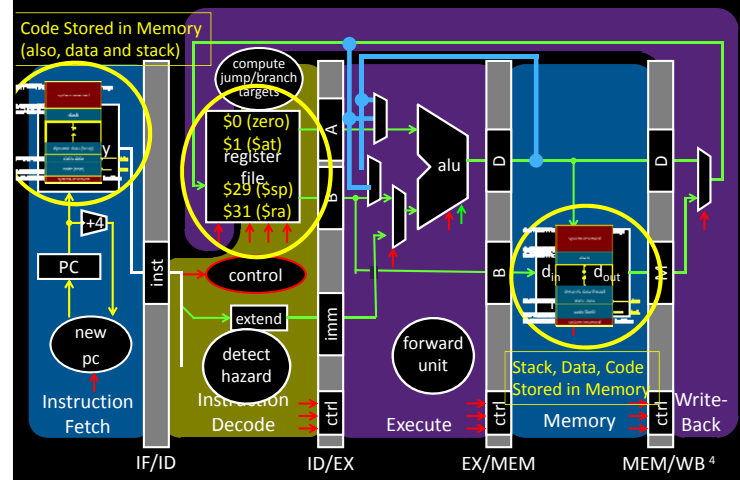
2

Agenda


- Memory hierarchy
- Cache
- Managing cache
- Write on cache
- Cache conscious programming

3

Big Picture: Memory




CPU/Memory performance



How do we make the processor fast,
Given that memory is VEEERRRRYYY SLLOOOWWW!!

5

Memory Pyramid



6


Memory Hierarchy

Memory closer to processor

- small & fast
- stores active data

Memory further from processor

- big & slow
- stores inactive data



7

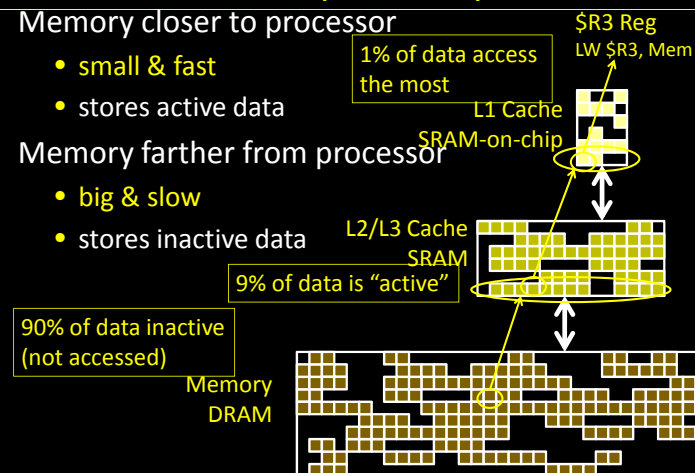
Memory Hierarchy

Memory closer to processor

- small & fast
- stores active data

Memory further from processor

- big & slow
- stores inactive data



8

Memory Hierarchy

Insight for Caches

If Mem[x] was accessed *recently*...

... then Mem[x] is likely to be accessed *soon*

- Exploit **temporal locality**:
 - Put recently accessed Mem[x] higher in memory hierarchy since it will likely be accessed again soon

... then Mem[x ± ε] is likely to be accessed *soon*

- Exploit **spatial locality**:
 - Put entire block containing Mem[x] and surrounding addresses higher in memory hierarchy since nearby address will likely be accessed

9

Memory Hierarchy

```
int n = 4;
int k[] = { 3, 14, 0, 10 };
int sum = 0;
```

```
int main(int ac, char **av) {
    for (int i = 0; i < n; i++) {
        sum += k[i];
        prints("\n");
    }
}
```

Temporal Locality

Spatial Locality

10

Memory Hierarchy


Memory closer to processor is fast but small

- Transfer whole **blocks** (cache lines):

4kb: disk ↔ ram

256b: ram ↔ L2

64b: L2 ↔ L1



11

Quiz

Which of the following statements are correct?

- The execution of an instruction always involves accessing a memory device.
- Over the last 30 years, the speed gap between memory and CPU is widening.
- To bridge the speed gap between memory and CPU, cache is used to store data permanently.
- Temporal locality means data items should be stored close to each other.
- The spatial locality exhibited by programs means cache hit rate can be improved by loading a block of data to the cache when a cache miss occurs.

12

Agenda

- Memory hierarchy
- **Cache**
- Managing cache
- Write on cache
- Cache conscious programming

13

Cache Lookups

Processor tries to access Mem[x]

Check: is block containing Mem[x] in the cache?

- Yes: **cache hit**
 - return requested data from cache line
- No: **cache miss**
 - read block from memory (or lower level cache)
 - (evict an existing cache line to make room)
 - place new block in cache
 - return requested data

14

Cache look up example

- Assume a cache has 4 lines (slots) and each line can hold one item (byte).
 - There are 4 slots where we can keep numbered (i.e. addressable) items.
- We want to keep in the 4 slots some items that have been used in the recent past.
- What process should we follow to make our cache work?

15

Example

address	Item	address	Item
		0	Geoff
		1	Joe
		2	Nora
		3	Allan
		4	Rick
		5	Gary
		6	Ben
		7	Chris
		8	Mitch
		9	Dan
		10	Sue
		11	Fred

Cache

List of Items

16

Example

Address	Item	Address	Item
10	Sue	0	Geoff
		1	Joe
		2	Nora
		3	Allan
		4	Rick
		5	Gary
		6	Ben
		7	Chris
		8	Mitch
		9	Dan
		10	Sue
		11	Fred

Cache

List of Items

17

Example

Address	Item	Address	Item
10	Sue	0	Geoff
3	Allan	1	Joe
		2	Nora
		3	Allan
		4	Rick
		5	Gary
		6	Ben
		7	Chris
		8	Mitch
		9	Dan
		10	Sue
		11	Fred

Cache

List of Items

18

Example

Address	Item	Address	Item
10	Sue	0	Geoff
3	Allan	1	Joe
7	Chris	2	Nora
		3	Allan
		4	Rick
		5	Gary
		6	Ben
		7	Chris
		8	Mitch
		9	Dan
		10	Sue
		11	Fred

Cache

List of Items

19

Example

Address	Item	Address	Item
10	Sue	0	Geoff
3	Allan	1	Joe
7	Chris	2	Nora
6	Ben	3	Allan
		4	Rick
		5	Gary
		6	Ben
		7	Chris
		8	Mitch
		9	Dan
		10	Sue
		11	Fred

Cache

Items are placed in the cache in the order that they are accessed

List of Items

20

Issues with our Cache

Searching through the cache for a given address is expensive.

- What is the time complexity of the search, in terms of the size of the cache?

21

Issues with our Cache

Is there a way to reduce the complexity (i.e. speed up the search)?

22

Direct-Mapped Caches

In our first attempt, any address from the list can go to any slot in the cache – thus there was a need to search *every* slot in the cache, when we looked for an address.

Now, we have a variant of the cache where an address can go to *just one* predefined slot in the cache.

- Address 0 will go to slot 0 (and only slot 0); address 1 will go to slot 1 (and only slot 1); etc;
- Address 4 will go to slot 0 (and only slot 0); address 5 will go to slot 1 (and only slot 1); etc;
- This is called a **direct-mapped** cache. In contrast, our first cache is called a **fully-associative** cache (where an address can go to any slot).

23

Direct-Mapped Caches

Address	Item	Address	Item
		0	Geoff
		1	Joe
		2	Nora
		3	Allan
		4	Rick
		5	Gary
		6	Ben
		7	Chris
		8	Mitch
		9	Dan
		10	Sue
		11	Fred

Cache

Which slot in the cache each address of the data will go to?

List of Items

24

Direct-Mapped Cache

Do you see any improvement in the search in direct-mapped caches?

25

Direct-Mapped Cache

Do you see any improvement in the search in direct-mapped caches?

```
where2look = myId % CacheSize // % is the mod operation
```

```
Check whether myAddr == cacheSlot[where2look].Addr
```

26

Direct-Mapped Cache

There is now only one slot where a particular address can be. We don't therefore need to search every slot.

We improve the complexity.

- Well, do we?

27

Direct-Mapped Cache

What price is this:

```
where2look = myId % CacheSize
```

How do we do this mod better?

28

Direct-mapped Caches

In hardware, we can replace the mod operation by a mask operation provided that the denominator is a power of 2.

11011010 AND 00000011

Direct-mapped caches in computers therefore have number of entries that is always a power of 2.

- There are many other things that are powers of 2 when it comes to computing hardware. Can you think of some others?

29

Quiz

Which of the following statements are correct?

- In a fully associative cache, a data item can be stored in any line (slot) of the cache.
- In a direct mapped cache, a data item must always be stored in the same cache line.
- Assume that a direct mapped cache has 64 slots and each slot holds one byte. Which slot is address 0x6798 mapped to?

30

Direct-Mapped Caches

Consider a direct-mapped cache of size 4. Each slot in the cache can have just one item (i.e. the line size is 1). The cache is empty to start with.

Work out if the following accesses to the given addresses are hits or misses. Each access is numbered with a sequence id for your convenience.

Sequence Id	1	2	3	4	5	6	7	8
Address	8	8	7	3	7	3	7	3
Hit/Miss								

31

Direct-Mapped Caches

Line #	Address	Item
0		
1		
2		
3		

Sequence Id	1	2	3	4	5	6	7	8
Address	8	8	7	3	7	3	7	3
Hit/Miss								

32

Direct-Mapped Caches

Line #	Address	Item
0	8	
1		
2		
3		

Sequence Id	1	2	3	4	5	6	7	8
Address	8	8	7	3	7	3	7	3
Hit/Miss	M							

33

Direct-Mapped Caches

Line #	Address	Item
0	8	
1		
2		
3		

Sequence Id	1	2	3	4	5	6	7	8
Address	8	8	7	3	7	3	7	3
Hit/Miss	M	H						

34

Direct-Mapped Caches

Line #	Address	Item
0	8	
1		
2		
3	7	

Sequence Id	1	2	3	4	5	6	7	8
Address	8	8	7	3	7	3	7	3
Hit/Miss	M	H	M					

35

Direct-Mapped Caches

Line #	Address	Item
0	8	
1		
2		
3	3	

Sequence Id	1	2	3	4	5	6	7	8
Address	8	8	7	3	7	3	7	3
Hit/Miss	M	H	M	M				

36

Direct-Mapped Caches

Line #	Address	Item
0	8	
1		
2		
3	7	

Sequence Id	1	2	3	4	5	6	7	8
Address	8	8	7	3	7	3	7	3
Hit/Miss	M	H	M	M	M			

37

Direct-Mapped Caches

Line #	Address	Item
0	8	
1		
2		
3	3	

Sequence Id	1	2	3	4	5	6	7	8
Address	8	8	7	3	7	3	7	3
Hit/Miss	M	H	M	M	M	M		

38

Direct-mapped Caches

Direct-mapped caches can lead to thrashing: two items go in and out all the time even though there is plenty of space available.

- Image the access order 3,7,3,7,3 in the previous example

Is there a way to get rid of thrashing?

- Fully-associative cache – our first cache where an address can go to any slot

Is there a way to reduce thrashing?

39

Set-Associative Caches

address	Item

Cache


We have two half-size direct-mapped caches side-by-side. Which slot in the cache each id of the data will go to?

address	Item
0	Geoff
1	Joe
2	Nora
3	Allan
4	Rick
5	Gary
6	Ben
7	Chris
8	Mitch
9	Dan
10	Sue
11	Fred

List of Items

40

Set-Associative Caches



2-way set associative cache

address	Item
0	Geoff
1	Joe
2	Nora
3	Allan
4	Rick
5	Gary
6	Ben
7	Chris
8	Mitch
9	Dan
10	Sue
11	Fred

List of Items

41

Set-Associative Caches


Consider a 2-way set associative cache of size 4. Each slot in the cache can have just one item (i.e. the line size is 1 item). The cache is empty to start with.

Work out if the following accesses to the given addresses are hits or misses. Each access is numbered with a sequence id for your convenience.

Sequence Id	1	2	3	4	5	6	7	8
Address	8	8	7	3	7	3	7	3
Hit/Miss								

42

Set-Associative Caches



Cache

We have 2 slots where each data id can go to. This reduces thrashing (does it?) but does not prevent it.

Id	Item
0	Geoff
1	Joe
2	Nora
3	Allan
4	Rick
5	Gary
6	Ben
7	Chris
8	Mitch
9	Dan
10	Sue
11	Fred

List of Items

43

Compromise

Set-associative cache

Like a direct-mapped cache

- Index into a location
- Fast

Like a fully-associative cache

- Can store multiple entries
 - decreases thrashing in cache
- Search in each set

44

Misses

Three types of misses

- **Compulsory**
 - The line is being referenced for the first time
- **Capacity**
 - Miss due to the cache not being big enough to hold the current data set. If the number of active lines is more than the cache can contain, capacity misses take place.
- **Conflict**
 - The required line was previously there in the cache, but was replaced by another line which happened to map to the same cache location.
 - Conflict misses take place because of limited or zero associativity when lines must be discarded in order to accommodate new lines which are mapped to the same line in the cache.
 - A miss occurs when the replaced line needs to be accessed again.

45

Misses

Q: How to reduce...

Compulsory Misses

- Unavoidable? The data was never in the cache...
- Prefetching!

Capacity Misses

- Buy more SRAM

Conflict Misses

- Use a more flexible cache design

46

Quiz

Which of the following statements are correct?

- Cache thrashing happens when all slots in the cache are occupied.
- Fully associative cache does not suffer from thrashing.
- In a 2-way set associative cache, an item in the memory can be mapped to 2 slots in the cache.
- Compared with a fully associative cache of the same size, the access speed of a 2-way set associative cache is faster.
- Compulsory miss can be avoided by increasing the size of a cache.
- Conflict miss can be avoided by using 2-way set associative cache.

47

Agenda

- Memory hierarchy
- Cache
- **Managing cache**
- Write on cache
- Cache conscious programming

48

Eviction

Which cache line should be evicted from the cache to make room for a new line?

- Direct-mapped
 - no choice, must evict line selected by index
- Associative caches
 - LRU: replace line that has not been used in the longest time

49

Replacing cache line: LRU

Consider a fully-associative cache of size 4. Each slot in the cache can have just one item (i.e. the line size is 1 item). The cache is empty to start with.

The cache uses an LRU replacement policy: every slot has a counter; every time a slot is accessed, a global counter is incremented and the value is stored in the slot counter; the slot with the lowest counter value is chosen for replacement.

Sequence Id	1	2	3	4	5	6	7	8	9	10
Address	7	8	7	8	7	3	7	3	4	2
Hit/Miss										

Fully-Associative Cache: LRU

Line #	Address	Item	Local Counter
0			
1			
2			
3			

Global Counter
Value

Sequence Id	1	2	3	4	5	6	7	8	9	10
Address	7	8	7	8	7	3	7	3	4	2
Hit/Miss										

Line #	Address	Item	Local Counter
0	7		1
1			
2			
3			

Global Counter	
Value	1

Sequence Id	1	2	3	4	5	6	7	8	9	10
Address	7	8	7	8	7	3	7	3	4	2
Hit/Miss	M									

Line #	Address	Item	Local Counter
0	7		1
1	8		2
2			
3			

Global Counter	
Value	2

Sequence Id	1	2	3	4	5	6	7	8	9	10
Address	7	8	7	8	7	3	7	3	4	2
Hit/Miss	M	M								

Line #	Address	Item	Local Counter
0	7		3
1	8		2
2			
3			

Global Counter	
Value	3

Sequence Id	1	2	3	4	5	6	7	8	9	10
Address	7	8	7	8	7	3	7	3	4	2
Hit/Miss	M	M	H							

Line #	Address	Item	Local Counter
0	7		3
1	8		4
2			
3			

Global Counter	
Value	4

Sequence Id	1	2	3	4	5	6	7	8	9	10
Address	7	8	7	8	7	3	7	3	4	2
Hit/Miss	M	M	H	H						

Line #	Address	Item	Local Counter
0	7		5
1	8		4
2			
3			

Global Counter	
Value	5

Sequence Id	1	2	3	4	5	6	7	8	9	10
Address	7	8	7	8	7	3	7	3	4	2
Hit/Miss	M	M	H	H	H					

Line #	Address	Item	Local Counter
0	7		5
1	8		4
2	3		6
3			

Global Counter	
Value	6

Sequence Id	1	2	3	4	5	6	7	8	9	10
Address	7	8	7	8	7	3	7	3	4	2
Hit/Miss	M	M	H	H	H	M				

Line #	Address	Item	Local Counter
0	7		7
1	8		4
2	3		6
3			

Global Counter	
Value	7

Sequence Id	1	2	3	4	5	6	7	8	9	10
Address	7	8	7	8	7	3	7	3	4	2
Hit/Miss	M	M	H	H	H	M	H			

Line #	Address	Item	Local Counter
0	7		7
1	8		4
2	3		8
3			

Global Counter	
Value	8

Sequence Id	1	2	3	4	5	6	7	8	9	10
Address	7	8	7	8	7	3	7	3	4	2
Hit/Miss	M	M	H	H	H	M	H	H		

Line #	Address	Item	Local Counter
0	7		7
1	8		4
2	3		8
3	4		9

Global Counter	
Value	9

Sequence Id	1	2	3	4	5	6	7	8	9	10
Address	7	8	7	8	7	3	7	3	4	2
Hit/Miss	M	M	H	H	H	M	H	H	M	

Line #	Address	Item	Local Counter
0	7		7
1	8 2		10
2	3		8
3	4		9

Global Counter	
Value	10

Sequence Id	1	2	3	4	5	6	7	8	9	10
Address	7	8	7	8	7	3	7	3	4	2
Hit/Miss	M	M	H	H	H	M	H	H	M	M

Memory Address

The memory is viewed as a large, single-dimension array.

A memory *address* is an index into the array

Byte addressing means that the index points to a byte of memory.

0	8 bits of data
1	8 bits of data
2	8 bits of data
...	8 bits of data
...	8 bits of data

Memory Address

Bytes are nice, but most data items use larger "words"

- Words are the smallest units of data we get out of/into memory

A word is 32 bits or 4 bytes in a 32-bit CPU.

- 2^{32} bytes with byte addresses from 0 to $2^{32}-1$
- 2^{30} words with word addresses 0, 4, 8, ... $2^{32}-4$

Similarly, in a 16-bit CPU (e.g. LC-3), a word is 16 bits

0	8 bits of data	0	32 bits of data	0	16 bits of data
1	8 bits of data	4	32 bits of data	2	16 bits of data
2	8 bits of data	8	32 bits of data	4	16 bits of data
3	8 bits of data	12	32 bits of data	6	16 bits of data
...	8 bits of data	...	32 bits of data	...	16 bits of data

Byte Address vs. Word Address


Byte Address is the fundamental address in a memory system, and is the address of a byte.

Word Address is the address of a word.

- Not all byte addresses are valid word addresses.

0	8 bits of data	0	32 bits of data	0	16 bits of data
1	8 bits of data	4	32 bits of data	2	16 bits of data
2	8 bits of data	8	32 bits of data	4	16 bits of data
3	8 bits of data	12	32 bits of data	6	16 bits of data
...	8 bits of data	...	32 bits of data	...	16 bits of data

- A cache is organised as a collection of lines.
- Each line consists of one or several bytes.




- A cache with 2 lines where each line has four words and each word has four bytes. How do we map a memory location to a cell in the cache (assume the cache is a direct mapped cache)?


65

- The total number of bytes in the cache is
 - num of lines \times num of words in a line \times num of bytes in a word
 - $2 \times 4 \times 4 = 32$
- To determine the cell in the cache, we carry out a MOD operation
 - address MOD 32
 - look at the last five bits of the address

66



67



68

- In the 5 least significant bits of the address, the left most bit determines the cache line and the rest of the bits decide the position of a byte in a line.
- In general, if there are 2^n cache lines and each line has 2^m bytes,
 - There are 2^{m+n} bytes in the cache.
 - The $m+n$ least significant bits of a memory address are used to determine where a byte should be stored in a cache.
 - Among the $m+n$ bits, the left-most n bits decides which cache line a byte is stored, and the other m bits indicate position of the byte with a cache line.

69


Mapping memory address to cache

Generally speaking, for direct mapped and set associative cache, when deciding which location in the cache a memory location should be stored, the memory address is partitioned into three fields.

- The offset field determines the position in a cache line.
- The index field decides the cache line.
- The tag field is used to distinguish different memory locations that are mapped to the same cache location.



70



71

- If each cache line only has one byte, the offset field is not needed.
 - Each line only has one location.
 - The address consists of two fields, i.e. tag and index

72

Tag

- For a direct mapped cache of size 32 and each slot only holds one item, where the data items at memory addresses 0x0 and 0x20 should be stored?

73


Tag

- For a direct mapped cache of size 32 and each slot only holds one item, where the data items at memory addresses 0x0 and 0x20 should be stored?
- They are both mapped to line 0.
- How can we know whether the value stored in line 0 is the data at address 0x0 or 0x20?

74

Tag

- Each cache line has a tag field that stores the value in the tag field of the address.



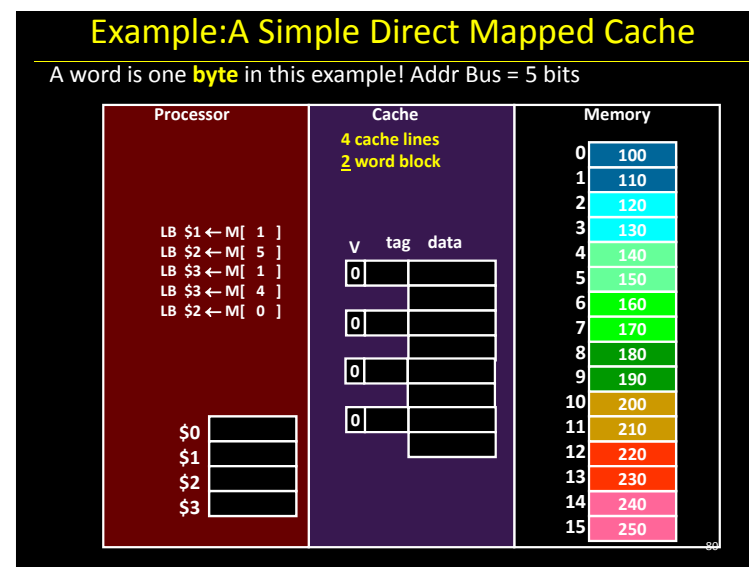
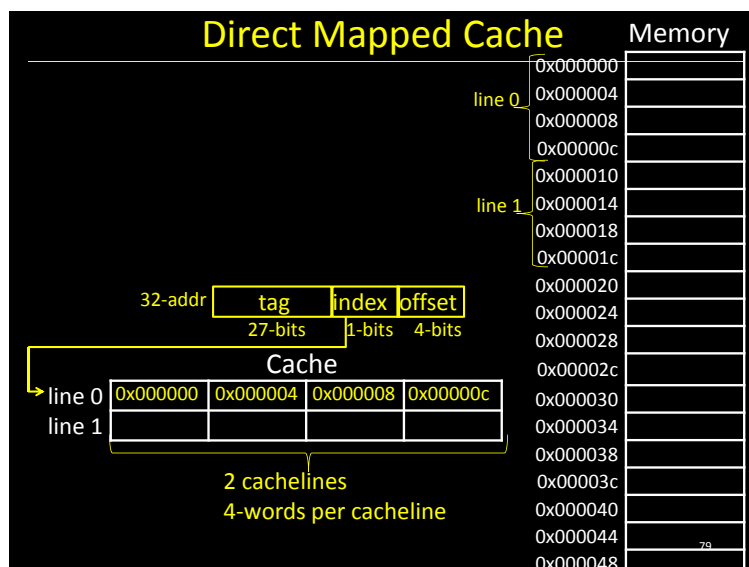
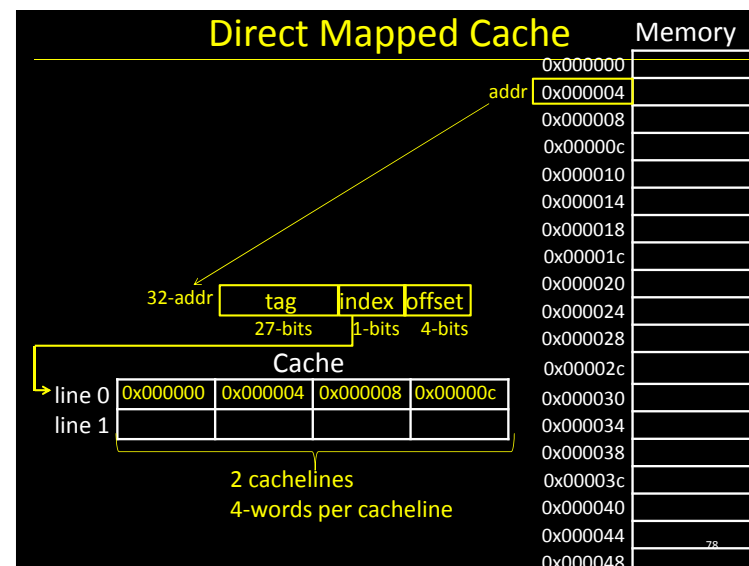
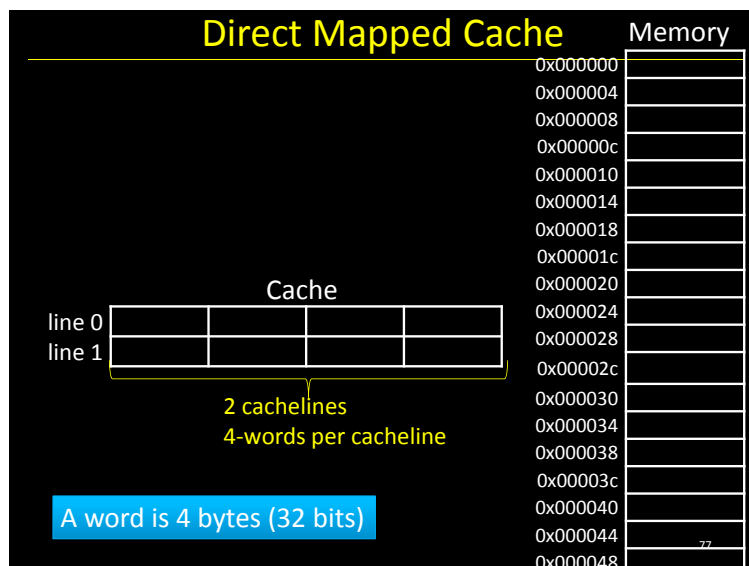
75

Valid bit

The valid bit indicates whether the data in a cache line is valid.

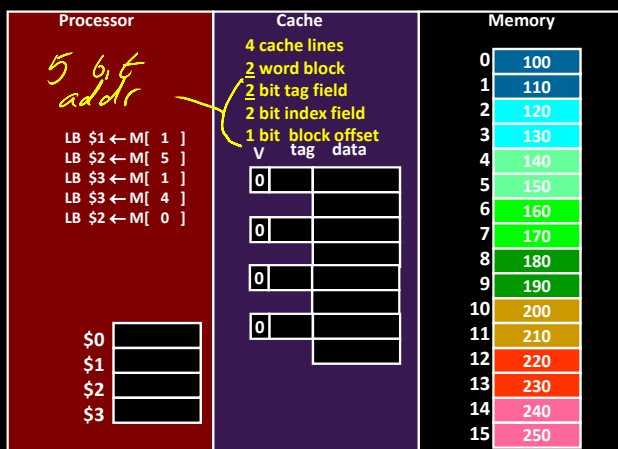
- Many programs share the cache

76



Example: A Simple Direct Mapped Cache

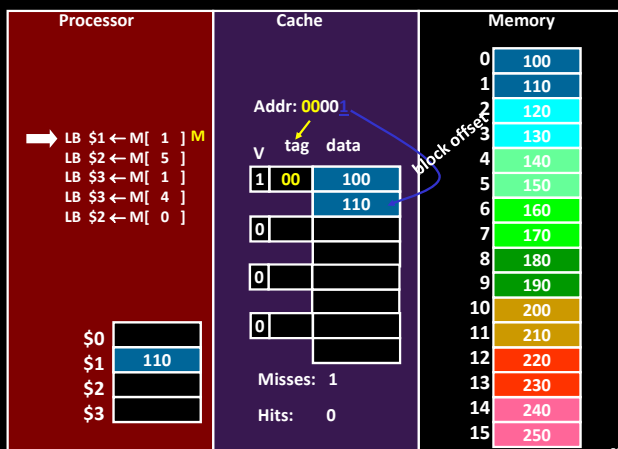
Using **byte addresses** in this example! Addr Bus = 5 bits



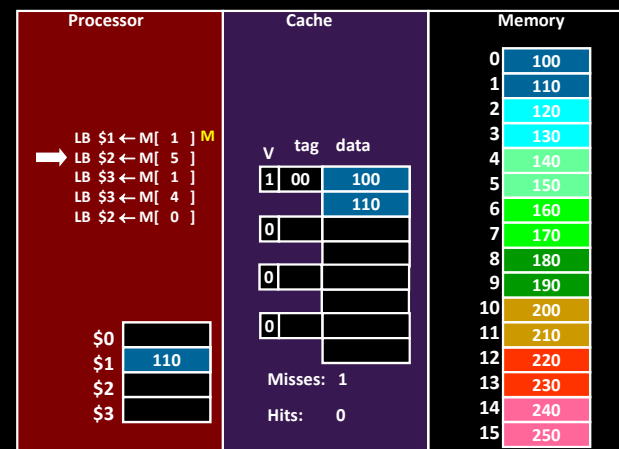
1st Access

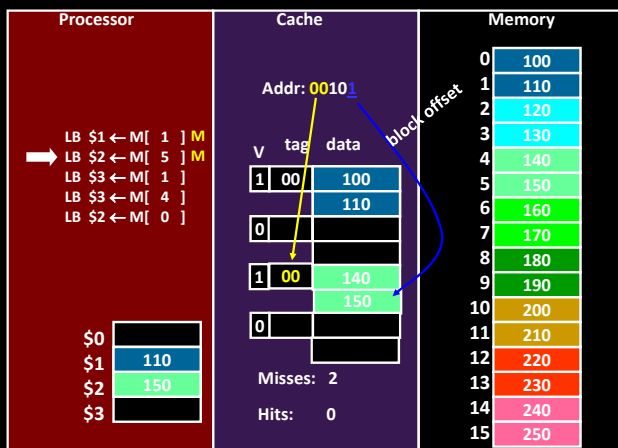
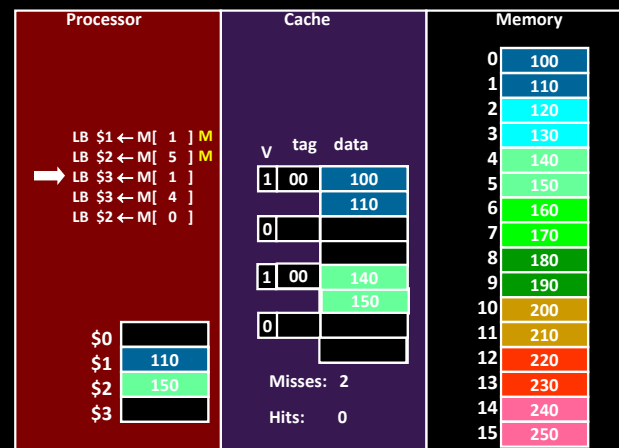
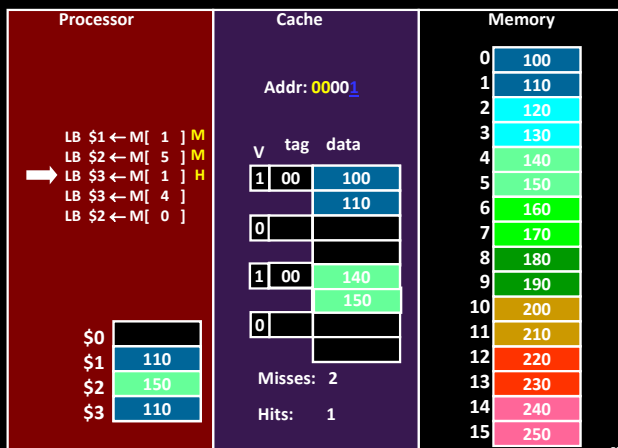
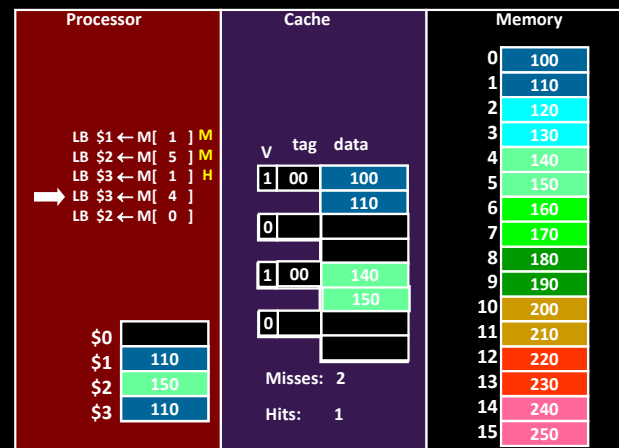


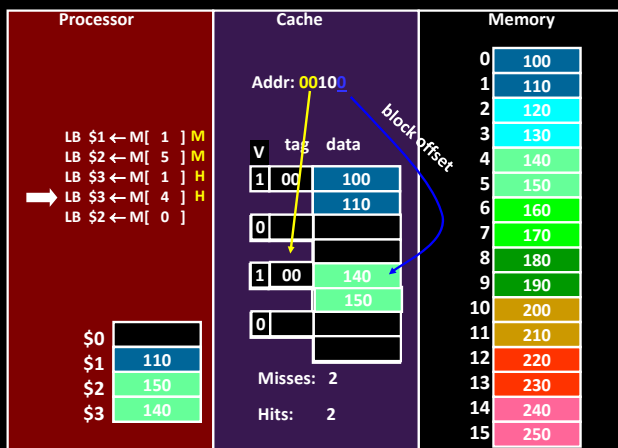
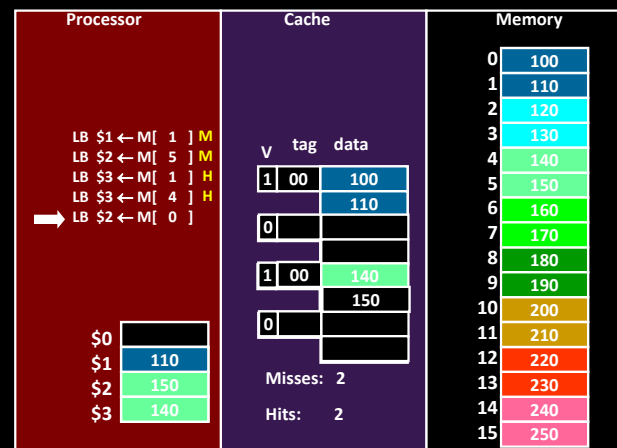
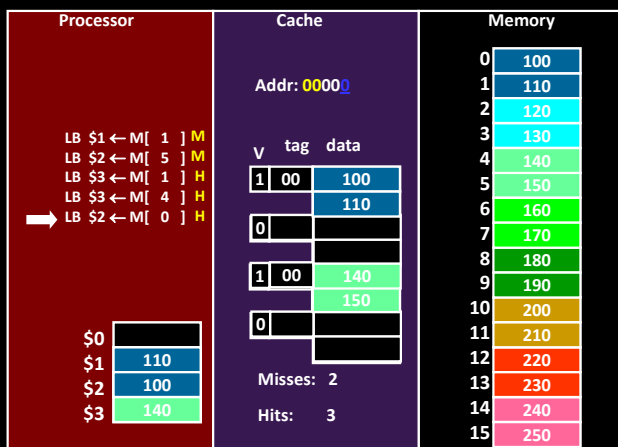
1st Access



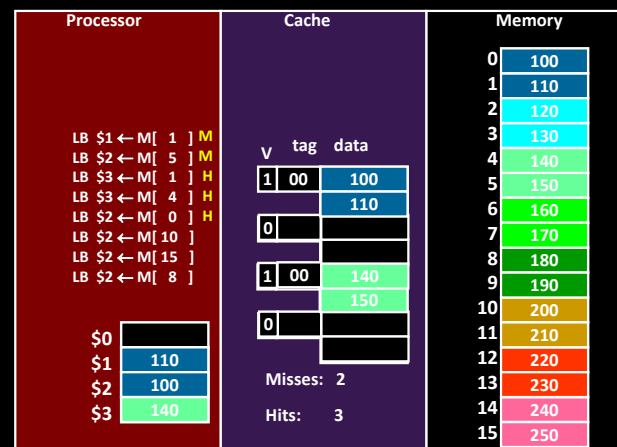
2nd Access



2nd Access3rd Access3rd Access4th Access

4th Access5th Access5th Access

Exercise



Quiz

Which of the following statements are correct?

- LRU is used to select the line to be evicted in direct mapped cache.
- When using a global and local counter to implement LRU cache eviction scheme, the local counter of a slot is incremented when a new value is written to the slot.
- In modern computer architecture, bytes are the smallest units of data that we get out of/into memory in one read/write operation.
- The size of a word is always 32 bits (i.e. four bytes).
- Assume that (a) a cache has 8 lines, (b) each line has 2 words, and (c) each word is 32 bits. The cache has 64 bytes.
- The tag field of a cache line stores the address of a memory location.
- The valid bit of a cache line indicates whether the contents of the line correspond to values stored in the memory.
- A direct mapped cache has 4 cache lines. Each cache line consists of 2 words, and each word is one byte. The address bus consists of 6 bits. What is the number of bits for the tag, index and offset fields respectively?

93


Agenda

- Memory hierarchy
- Cache
- Managing cache
- **Write on cache**
- Cache conscious programming

94

Write Policies

Q: How to write data?



If data is already in the cache...

Write-Through

- writes go to main memory and cache

Write-Back

- CPU writes only to cache
- cache writes to main memory later (when block is evicted)

95

Write Policies

- If it is not in the cache?
 - Allocate the line (put it in the cache)
(**write allocate policy**)
 - Write it directly to memory without allocation in the cache
(**no write allocate policy**)

96

write-through

A cache with a write-through and write-allocate policy

- reads an entire block (cache line) from memory on a cache miss
- writes only the updated item to memory for a store
- evictions do not need to write to memory

97

write-back

NOT write all stores immediately to memory

- Keep the most current copy in cache, and update memory when that data is *evicted* (**write-back policy**)

98

Write-Back Meta-Data

V	D	Tag	Byte 1	Byte 2	... Byte N

V = 1 means the line has valid data

D = 1 means the bytes are newer than main memory

When allocating line:

- Set V = 1, D = 0, fill in Tag and Data

When writing line:

- Set D = 1

When evicting line:

- If D = 0: just set V = 0
- If D = 1: write-back Data, then set D = 0, V = 0

99

Write-Back

Each miss reads a block

Each evicted dirty cache line writes a block

No write is needed if the cache line is not dirty

100

Performance Tradeoffs

In a cache hit, write-through is slower than write-back on writes.

In a cache miss, write-back might be slower than write-through if a dirty line has to be evicted.

101

Quiz

Which of the following statements are correct?

- The write-through policy modifies both the contents of the cache and the memory in a write operation.
- If a write operation has a cache hit, the write-back policy would delay the write operation on the cache until the cache line is evicted.
- When a cache line is evicted, the write-back policy always writes a complete cache line to the memory.
- Compared with the write-through policy, the write-back policy always gives better performance to the cache.

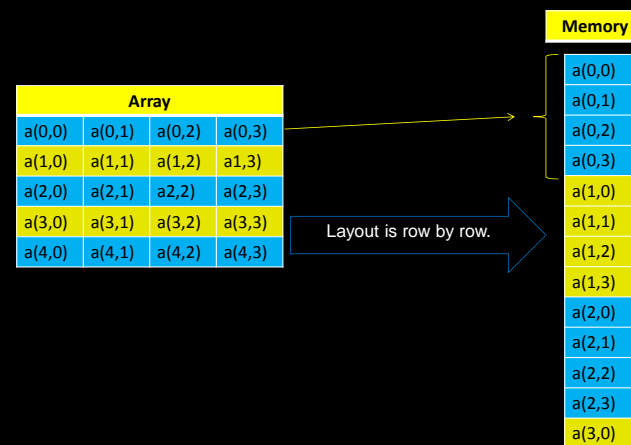
102

Agenda

- Memory hierarchy
- Cache
- Managing cache
- Write on cache
- **Cache conscious programming**

103

C's Array Layout in Memory



104

Two ways to sum a 2-D array

Process by row by row

```
#define N 10000
int a[N][N];
int main() {
    int sum=0, i, j;
    for (i=0; i<N; i++)
        for (j=0; j<N; j++)
            sum += a[i][j];
}
```

a(0,0), a(0,1), a(0,2), a(0,3), ...

a(1,0), a(1,1), a(1,2), a(1,3), ...

a(2,0), a(2,1), ...

Access order = memory layout order

Memory

a(0,0)
a(0,1)
a(0,2)
a(0,3)
a(1,0)
a(1,1)
a(1,2)
a(1,3)
a(2,0)
a(2,1)
a(2,2)
a(2,3)
a(3,0)

105

Two ways to sum a 2-D array

Process by column by column

```
#define N 10000
int a[N][N];
int main() {
    int sum=0, i, j;
    for (j=0; j<N; j++)
        for (i=0; i<N; i++)
            sum += a[i][j];
}
```

a(0,0), a(1,0), a(2,0), a(3,0), ...

a(0,1), a(1,1), a(2,1), a(3,1), ...

a(0,2), a(1,2), ...

Access order != memory layout order

Memory

a(0,0)
a(0,1)
a(0,2)
a(0,3)
a(1,0)
a(1,1)
a(1,2)
a(1,3)
a(2,0)
a(2,1)
a(2,2)
a(2,3)
a(3,0)

106

Process by row

Assume a is a 4x4 array

A cache line can hold 4 array elements

The cache has a single line

when i=0

Cache			
a(0,0)	a(0,1)	a(0,2)	a(0,3)

a(0,0), a(0,1), a(0,2), a(0,3),

a(1,0), a(1,1), a(1,2), a(1,3),

a(2,0), a(2,1), ...

Access order = memory layout order

one miss for each row

4 misses in total

Memory

a(0,0)
a(0,1)
a(0,2)
a(0,3)
a(1,0)
a(1,1)
a(1,2)
a(1,3)
a(2,0)
a(2,1)
a(2,2)
a(2,3)
a(3,0)

107

Process by column

Assume a is a 4x4 array

A cache line can hold 4 array elements

The cache has a single line

when j=0 and i=0

Cache			
a(0,0)	a(0,1)	a(0,2)	a(0,3)

a(0,0), a(1,0), a(2,0), a(3,0), ...

a(0,1), a(1,1), a(2,1), a(3,1), ...

a(0,2), a(1,2), ...

Access order != memory layout order

How many misses for each column?

How many misses in total?

Memory

a(0,0)
a(0,1)
a(0,2)
a(0,3)
a(1,0)
a(1,1)
a(1,2)
a(1,3)
a(2,0)
a(2,1)
a(2,2)
a(2,3)
a(3,0)

108

Download and time the two example programs.

- `time ./cache-r`
- `time ./cache-w`

109