

# Multidimensional Data

---

- ❑ Similarity search in high-dimensional spaces
- ❑ Motivating applications
- ❑ Indexing issues and the curse of dimensionality
- ❑ Indexing based on dimensionality reduction
- ❑ Indexing based on compression
- ❑ Indexing metric spaces
- ❑ Subsequence matching in time-series


# Similarity search

---

- ❑ Record data representation
- ❑ Given an object (record)  $o$  find
  - similar objects to  $o$  according to some upper distance bound  $\varepsilon$
  - $k$ -most similar objects to  $o$
- ❑ Applications
  - multimedia search-by-example
  - data mining (clustering, NN-based classification)
- ❑ Issues
  - bad performance of spatial indexes in high-dimensional spaces
  - curse of dimensionality

# Application: Multimedia Data


---



- ❑ Objective: query and analyze vast amounts of multimedia data (e.g., images)
- ❑ *Content-based Image Retrieval*:
  - index and retrieve images based on their *visual contents* (e.g. color distribution)
- ❑ Methodology:
  - From each image extract a fixed set of *features* (e.g. color features)
  - Represent images as *feature vectors*
  - Index and query feature vectors instead of images

# Color Features

- ❑ To represent the color of an image compactly, a *color histogram* is used. Colors are partitioned to  $k$  groups according to their similarity and the *percentage* of each group in the image is measured.
- ❑ Images are transformed to  $k$ -dimensional points and a distance metric (e.g., Euclidean distance) is used to measure the similarity between them.



# Distance Metrics in a Multidimensional Space

---

□ Given two n-dimensional points

- $p = p_1 \dots p_n$

- $q = q_1 \dots q_n$

□ their Euclidean distance is defined as:

$$L_2(p, q) \equiv \sqrt{\sum_{i=1}^n (p_i - q_i)^2}$$

□ also

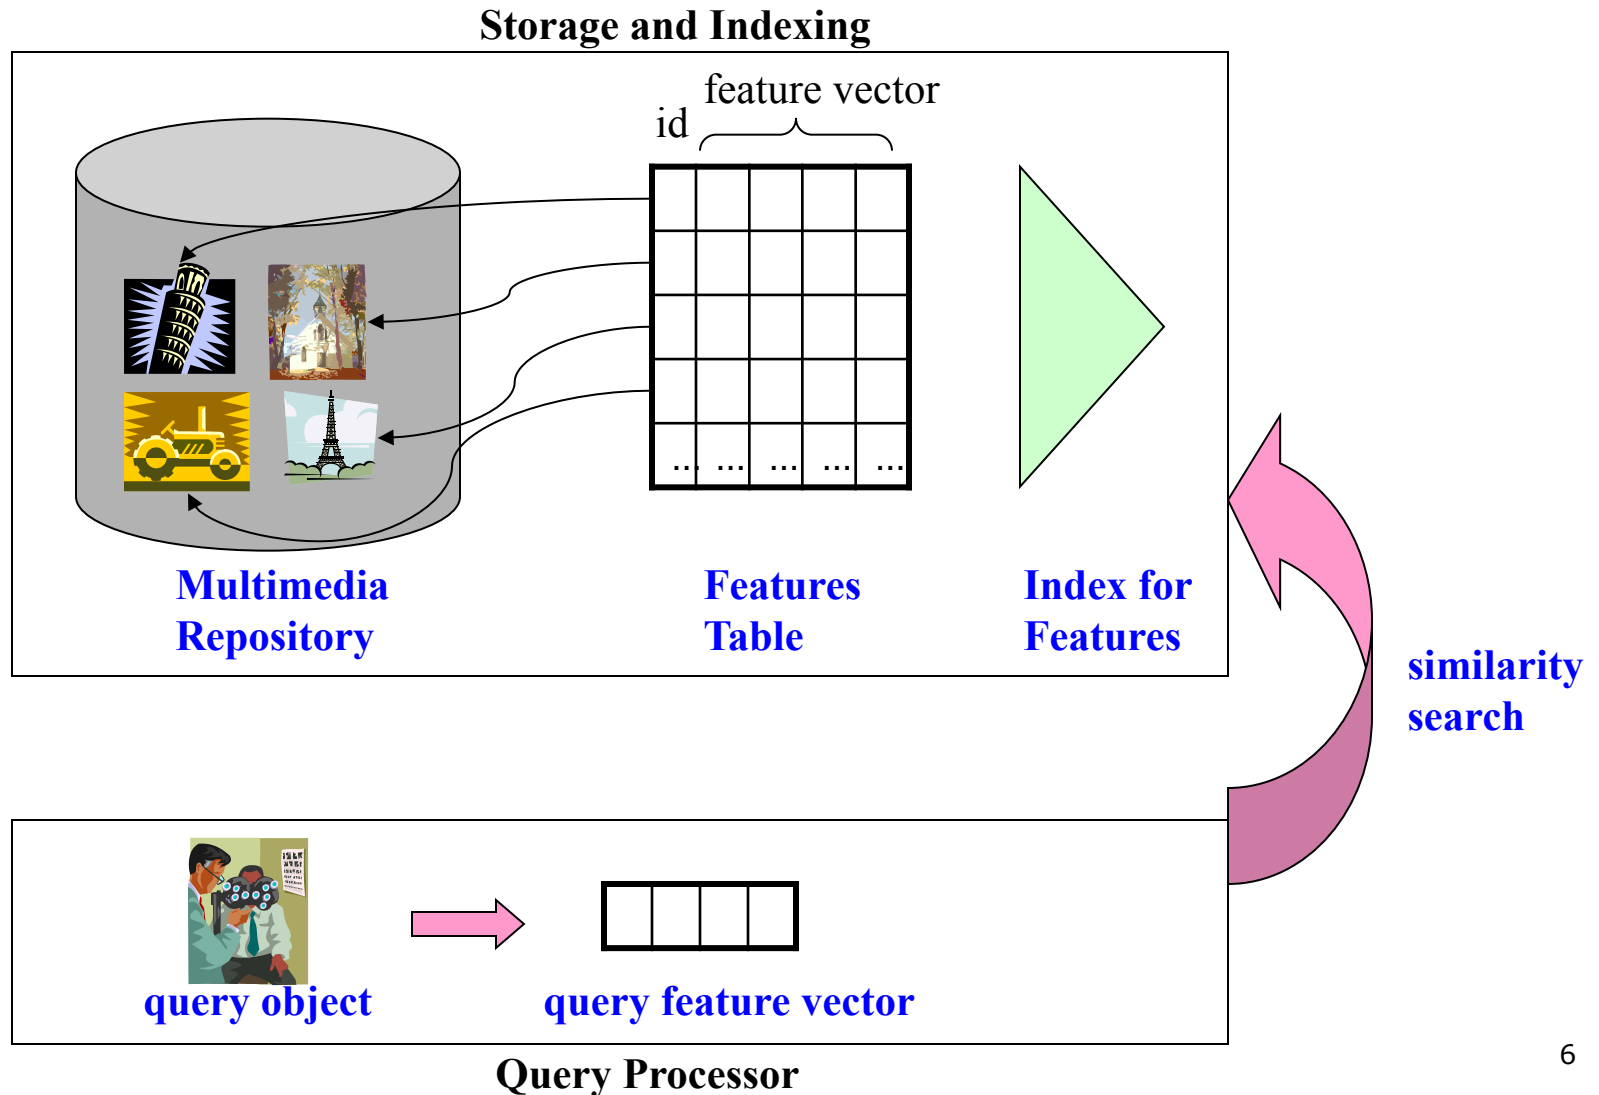
- Manhattan (city block) distance

$$L_1(p, q) \equiv \sum_{i=1}^n |p_i - q_i|$$


- Max (supremum) distance

$$L_\infty(p, q) \equiv \max_{i=1}^n |p_i - q_i|$$

# Architecture of a Multimedia Database



# Application: Time-series Data



- A time-series is a sequential collection of values or events over time.
- Time series data are found in everywhere, e.g., stock market values, sensor indications, cardiograms.


real-valued time sequence  
(e.g. stock prices)



event sequence  
(e.g., human activities during a day)





AABBBSSBABAGGTSAK

# Queries and Analysis Tasks on Time-series Data



- find the most similar sequence to a query sequence  $q$


1: Whole Matching



$C_6$  is the best match.





# Queries and Analysis Tasks on Time-series Data



- find (approximate/exact) occurrences of a query subsequence  $q$  in a long sequence  $T$ .


## 2: Subsequence Matching



# Queries and Analysis Tasks on Time-series Data




- Classification and Clustering (same are found in multimedia databases)



- Discovery of rules and trends
  - If stock A moves up for 10 days in a row then it will move down the next 5 days with high probability

# Problem: Indexing feature vectors or time sequences for fast similarity search

- **Possible solution:** represent each vector as a point in the multi-dimensional feature space, index them by an R-tree and use spatial query methods



# Problem!

---

- ❑ The R-tree does not scale well for many dimensions. Somewhere above 6-20 dimensions search using the R-tree is slower compared to linear scan
  - large MBRs, a lot of empty space
  - not all dimensions are used for splitting
    - ❑ a query point is inside ALL MBRs in most dimensions!
- ❑ Feature vectors and time-series are long (hundreds of dimensions)
- ❑ Distances between objects become meaningless even with a few noise dimensions (dimensionality curse)


# Possible Remedies

---

- The problem can be alleviated by:
  - dimensionality reduction and application of multi-step search algorithms
  - data compression and linear scan
  - indexing the metric distance space
  - approximate search

# Dimensionality Reduction

- ❑ In many cases the *embedded dimensionality* of a search problem is much lower than the actual dimensionality
- ❑ Some methods *apply transformations* on the data and approximate them with *low-dimensional vectors*
- ❑ The aim is to reduce dimensionality and at the same time maintain the data characteristics



# GEMINI GEneric Multimedia INdexIng

---

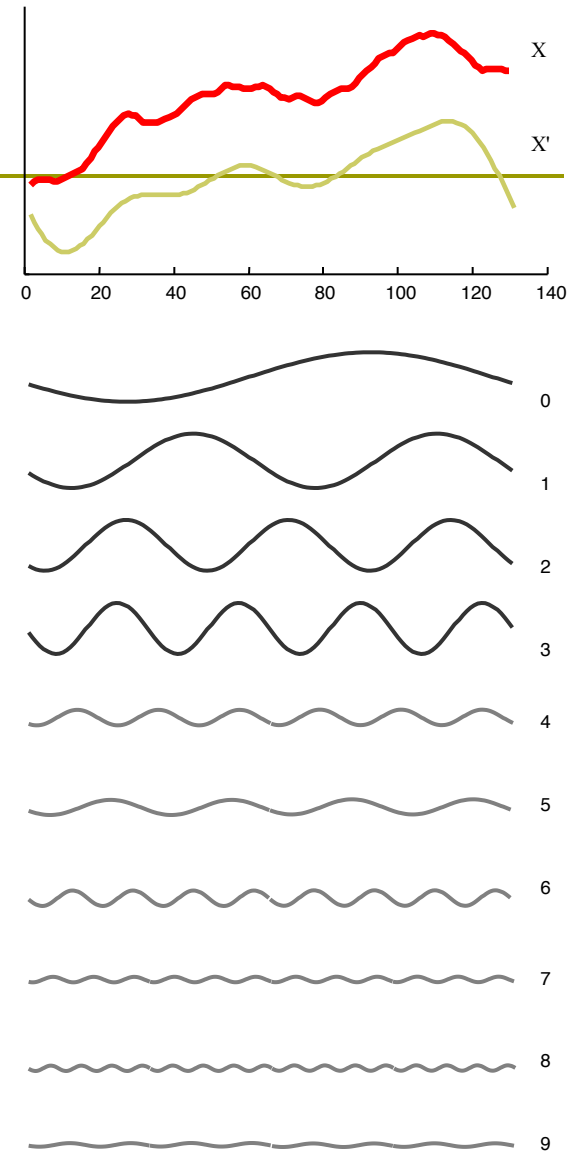
- ❑ Establish a distance metric from a domain expert.
  - e.g., Euclidean distance
- ❑ Produce a **dimensionality reduction technique** that reduces the dimensionality of the data from  $n$  to  $N$ , where  $N$  can be efficiently handled by a spatial access method (e.g., R-tree)
- ❑ Produce a distance measure defined on the  $N$  dimensional representation of the data, and **prove** that it obeys
  - $D'(A', B') \leq D(A, B)$
  - the above is called the **lower bounding** property
- ❑ Plug into an off-the-shelf spatial-access-method (e.g., R-tree).

# Dimensionality Reduction Example: Discrete Fourier Transform (DFT)

- How to represent a time-series (or a color histogram) using only  $n$  numbers (in a  $n$ -dimensional space)?
- Basic Idea: Represent the time series as a linear combination of sines and cosines, but keep **only the first  $n/2$  coefficients**.

Why  $n/2$  coefficients? Because each wave requires 2 numbers, for the phase ( $w$ ) and amplitude ( $A, B$ ).

$$C(t) = \sum_{k=1}^n (A_k \cos(2\pi w_k t) + B_k \sin(2\pi w_k t))$$





# Other Dimensionality Reduction Techniques

---

- ❑ Other popular transformations include:
  - **Discrete Wavelet Transform**. The sequence is transformed to a linear combination of Wavelet basis functions, but only the largest  $n$  coefficients are used.
  - **Singular Value Decomposition**. Similarly, the sequence is transformed to a linear combination of eigenwaves, and only the first  $n$  coefficients are used.
  - **Piecewise approximations** are also used for time-series data.

# Two-step processing of range similarity queries

---

## Range Similarity Queries

### □ Given:

- A database  $S$  of feature vectors (or time sequences)
- A distance function  $D(p_1, p_2)$  that computes the dissimilarity between vectors  $p_1$  and  $p_2$
- A query vector  $q$ , a distance threshold  $\varepsilon$

### □ Find:

- All vectors  $p$  in  $S$ , such that  $D(p, q) \leq \varepsilon$

# Two-step processing of range similarity queries

---

## □ Methodology:

- Each vector  $p$  in  $S$  is transformed to a low-dimensional vector  $p'$
- An index  $R$  for low-dimensional transformed vectors (e.g., R-tree) is used to for all **transformed vectors**  $p'$ .
- We define a distance function (e.g., Euclidean distance)  $D'(p', q')$  for the **transformed** vectors, such that
$$D'(p', q') \leq D(p, q) \text{ (lower bounding property)}$$

# Two-step processing of range similarity queries

---

## □ Step 1:


- convert  $q$  to  $q'$  using the same dimensionality reduction technique
- apply an **R-tree range search** to find fast a  $S' \subseteq S$ , such that for all  $p'$  in  $S'$ ,  $D'(p', q') \leq \epsilon$

## □ Step 2:


- for each  $p'$  in  $S'$ , get high-dimensional vector  $p$ , compute  $D(p, q)$  and if  $D(p, q) \leq \epsilon$  add it to the response set

# Two-step processing of range similarity queries

high-dimensional space



low-dimensional space



Q1: Will this method miss any results?

Q2: Will this method compute incorrect results?

# Two-step processing of nearest neighbor similarity queries

---

## Nearest Neighbor Similarity Queries

### □ Given:

- A database  $S$  of multimedia feature vectors (or time sequences).
- A distance function  $D(p_1, p_2)$  that computes the dissimilarity between vectors  $p_1$  and  $p_2$ .
- A query vector  $q$

### □ Find:

- The **most similar** vector to  $q$  in  $S$ .
- Or else,  $p \in S$  such that  $\forall s \in S, D(p, q) \leq D(s, q)$

# Two-step processing of nearest neighbor similarity queries

---

## □ Methodology:

- Each vector  $p$  in  $S$  is transformed to a low-dimensional vector  $p'$
- An index  $R$  for low-dimensional transformed vectors (e.g., R-tree) is used to for all **transformed vectors**  $p'$ .
- We define a distance function (e.g., Euclidean distance)  $D'(p',q')$  for the **transformed** vectors, such that
$$D'(p',q') \leq D(p,q) \text{ (lower bounding property)}$$

# Two-step processing of nearest neighbor similarity queries

---

## □ Step 1:

- convert  $q$  to  $q'$  using the same dimensionality reduction technique
- apply an **R-tree nearest neighbor** search to find fast the nearest  $p'$  to  $q'$ .
- Let  $p \in S$  be the corresponding high-dim vector to  $p'$ . Compute  $D(q, p)$ . Apply an **R-tree range search** to find fast a  $S' \subseteq S$ , such that for all points  $s'$  in  $S'$ ,  $D'(s', q') \leq D(q, p)$


## □ Step 2:

- for each  $s'$  in  $S'$  compute  $D(s, q)$  and return the one with the smallest  $D(s, q)$




# Two-step processing of nearest neighbor similarity queries

high-dimensional space



transformation space




- Q: Will this method compute the correct result? why?

# Multi-step processing of nearest neighbor similarity queries


---

- Convert  $q$  to  $q'$  using the same dimensionality reduction technique.
  - $NN = \text{NULL}; D(q, NN) = \infty;$
  - Repeat:
    - apply an **incremental R-tree nearest neighbor** search to find fast the next nearest  $p'$  to  $q'$ .
    - If  $D'(q', p') < D(q, NN)$  compute actual  $D(q, p)$ .
      - If  $D(q, p) < D(q, NN)$  then  $NN = p$
  - Until  $D'(q', p') \geq D(q, NN)$
- 
- **Q: Will this method better than two-step processing?**


# Example of multi-step processing




1. Get 1<sup>st</sup> NN ( $a'$ )



2. Compute  $a$ 's distance from  $q$ . Set  $cur_{NN}=a$ .





3. Get 2<sup>nd</sup> NN ( $a'$ ).  $D'(q',b') < D(q,a)$ ,  
so goto step 4



4. Compute  $b$ 's distance from  $q$ . Since  $D(q,b) < D(q,a)$ ,  
Set  $cur_{NN}=b$ .

# Example of multi-step processing



5. Get 3<sup>rd</sup> NN ( $c'$ ).  $D'(q', c') > D(q, b)$ ,  
so terminate and report  $b$  as NN.

# Compression-based similarity search

---

## □ Motivation:

- In very high dimensional spaces dimensionality reduction is hard to apply
  - Should examine **multiple possible dimension-sets** to potentially reduce
- Sometimes we have to resort to **linear scan**
  - Expensive because **the entire** (large) set of feature vectors have to be scanned and for each of them we need an (expensive) **distance computation**

# Compression-based similarity search


---

## □ Idea:

- Partition the space by a **grid**
- Approximate each vector by a **bitstring** that designates the **partition** where it belongs
- **Linearly scan all bitstrings** and use bounds to eliminate most of the objects
- Perform **exact distance computations** for the objects that survive the scan

# The Vector-Approximation File

## ▣ data preprocessing phase

















# The Vector-Approximation File

## □ similarity search (step 1)











- scan VA-file and for each bitstring  $b$ 
  - compute upper bound  $\text{dist}_u(q,b)$  to  $q$ 
    - keep track of smallest upper bound  $t$
  - compute lower bound  $\text{dist}_l(q,b)$  to  $q$
  - if  $\text{dist}_l(q,b) \geq t$  then filter out  $b$
  - else put  $b.\text{obj}$  to candidates set

original data

	00	01	10	11
00	 			
01				
10			 	
11				



VA-file


	0000	$t =$ 
	0000	$t =$ 
	0001	$t =$ 
	0100	$t =$ 
	<del>1010</del>	$t =$ 
	...	




# The Vector-Approximation File

## □ similarity search (step 2)


- Sort non-filtered candidates by  $\text{dist}_t(q,b)$ , scan them and for each bitstring  $b$ 
  - compute lower bound  $\text{dist}_t(q,b)$  to  $q$
  - if  $\text{dist}_t(q,b) \geq t$  then filter out  $b$
  - else compute actual  $\text{dist}(q,b.\text{obj})$  and update current actual NN and  $t$



VA-file



original data




# Indexing metric spaces: the M-tree

---

## □ Idea:

- Instead of indexing the feature space of the objects, index them by their **metric distances**
- Group objects into index nodes hierarchically using only the distances between them
  - Applicable even if the objects have unknown attributes and only the distances between them are known
- Each node  $n$  has a **routing object**  $o$  and a **radius**  $r$  (stored in the entry pointing to  $n$ )
  - For every object  $o'$  in the sub-tree of  $n$  it should be  $\text{dist}(o, o') \leq r$
  - All data objects appear in leaf nodes

# The M-tree: example



[Figure taken from <http://itu-algorithms.github.io/events/2015/05/20/Zezulacourse>]


# The M-tree: queries

## □ Use triangular inequality to search

example:

find all objects with distance  
at most 1.7 from  $q$

- 1)  $\text{dist}(q, o_1) > 4.5 + 1.7$ , so  $o_1$  and its subtree are pruned
- 2)  $\text{dist}(q, o_2) \leq 6.9 + 1.7$ , so  $o_2$  is followed
- 3)  $\text{dist}(q, o_7) \leq 1.3 + 1.7$ , so  $o_7$  is followed
- 4)  $o_7$  and  $o_{11}$  are added to the result
- 5)  $\text{dist}(q, o_2) \leq 2.9 + 1.7$ , so  $o_2$  is followed
- 6)  $o_8$  is added to the result
- 7)  $\text{dist}(q, o_4) \leq 1.6 + 1.7$ , so  $o_4$  is followed
- 8) no more results



[Figure taken from <http://itu-algorithms.github.io/events/2015/05/20/Zezulacourse>]

# The M-tree: updates


---

## □ Insertions:

- “Search” the tree, by recursively following the routing object which is the **closest** to the new object  $o$
- If the leaf node fits  $o$ , insert it there, otherwise **split** the leaf node
  - Partition the objects of the leaf +  $o$  to two new leaf nodes with **two new routing objects** for them and replace the old routing object in the parent by the two new routing objects
    - New routing objects should minimize volume and overlap of the new leaf nodes they define


## □ Deletions: as in R-tree

# Subsequence Matching



- find (approximate/exact) occurrences of a query subsequence  $q$  in a long sequence  $T$ .

## 2: Subsequence Matching



# Subsequence Matching

---

## □ Problem:

- Given one or possibly more long real-valued sequences (i.e., time-series), develop an index for subsequence matching queries:
- Subsequence Matching: given a database of long sequences  $S$  and a query sequence  $q$ , find all subsequences  $s$  in  $S$ , such that  $D(s,q) \leq \varepsilon$ , where  $\varepsilon$  is a **distance threshold**.

# Subsequence Matching

---


- Determine a short **sliding window**  $w$
- Assume for the moment that every query  $q$  is of length  $w$ 
  - shorter queries than  $w$  are not interesting
  - we will discuss about longer ones soon
- Assume that  $D =$  Euclidean distance



# Subsequence Matching


---

- Each position of the window defines a subsequence of length  $w$
- Each subsequence can be transformed to a point in a low dimensional space
  - E.g., using DFT, WDT, etc.
- The Euclidean low-dimensional distance lower-bounds  $D$ .



# Subsequence Matching

- The low-dimensional vectors of consecutive subsequences define a **trail** in the low-dimensional space




- Each point on a trail is a subsequence
- we can build an R-tree for these points

# Subsequence Matching

---


- However, the number of points can be too large, resulting in a large (and slow index)



- Example: 10M elements in a long vector, result in about 10M points

# Subsequence Matching

- Idea: divide the trails and approximate them by MBRs (hyper-rectangles)




- Each indexed MBR unit stores
  - a) the id of the sequence
  - b)  $t_{\text{start}}, t_{\text{end}}$  of the sub-trail
- Example: the 10M sequence can be divided into 10000 trails (about 1000 points per trail), thus the index will now have 10000 entries (instead of 10M)

# Subsequence Matching


## □ Searching:

- Use the R-tree to find fast the sub-trails close to the query in the low-dimensional space
- Linear-scan these trails to discard false-alarms



# Subsequence Matching

- Searching for subsequences  $q$  longer than  $w$ :
  - divide  $q$  into  $p$  segments, each of length  $w$
  - for each segment  $q_{\text{seg}}$ , apply an  $\epsilon/\sqrt{p}$  range query to find candidate subsequences that are close to the segment.
  - Unify the results for all segmental queries and examine the corresponding subsequences to discard false alarms
- Lemma:
  - If two sequences  $s$  and  $q$  are within distance  $\epsilon$  from each other then at least one pair of segments  $s_i$  and  $q_i$  should be within distance  $\epsilon/\sqrt{p}$  from each other



# Subsequence Matching: Example

---

- $T = \{3, 3, 4, 3, 5, 6, 7, 7, 8, 9, 10, 11, 9, 8, 9, 10\}$
- $w = 3$
- Subsequences:  $(3, 3, 4), (3, 4, 3), (4, 3, 5), \dots$
- Assume no dim. reduction (already 3 dimensions only!)
- Consider query  $q = \{3, 4, 3, 5, 6, 5\}$  and  $\epsilon = 4$  ( $\epsilon^2 = 16$ )
- Step 1: break  $q$  into  $q_1 = \{3, 4, 3\}$  and  $q_2 = \{5, 6, 5\}$
- Step 2:
  - use the index to search for subsequences with sq. distance at most  $\epsilon^2/2$  from  $q_1$  and put their position in  $S_1$
  - use the index to search for subsequences with sq. distance at most  $\epsilon^2/2$  from  $q_2$  and put their position in  $S_2$
  - Merge  $S_1$  and  $S_2$  to a set of candidate positions  $P$  to examine
- Step 3:
  - Perform random accesses to positions in  $P$  to examine the candidate  $q$ -length subsequences.

# Subsequence Matching: Example

$T = \{ \overset{1}{3}, \overset{2}{3}, \overset{3}{4}, \overset{4}{3}, \overset{5}{5}, \overset{6}{6}, \overset{7}{7}, \overset{8}{7}, \overset{9}{8}, \overset{10}{9}, \overset{11}{10}, \overset{12}{11}, \overset{13}{9}, \overset{14}{8}, \overset{15}{9}, \overset{16}{10} \}$

$q1 = \{3, 4, 3\}$

$\epsilon^2/2 = 8$

$$\underbrace{d^2(q1, 1)} = 0 + 1^2 + 1^2 = 2 \leq \epsilon^2/2 \quad \checkmark$$

$$\underbrace{d^2(q1, 2)} = 0 + 0 + 0 = 0 \leq \epsilon^2/2 \quad \checkmark$$

$$\underbrace{d^2(q1, 3)} = 1^2 + 1^2 + 2^2 = 6 \leq \epsilon^2/2 \quad \checkmark$$

$$\underbrace{d^2(q1, 4)} = 0^2 + 1^2 + 3^2 = 10 > \epsilon^2/2 \quad \text{X}$$

$$\underbrace{d^2(q1, 5)} = 2^2 + 2^2 + 4^2 = 24 > \epsilon^2/2 \quad \text{X}$$

.... remaining positions have sq. distance  $> \epsilon^2/2$

$S1 = \{1, 2, 3\}$



# Subsequence Matching: Example

$T = \{ \overset{1}{3}, \overset{2}{3}, \overset{3}{4}, \overset{4}{3}, \overset{5}{5}, \overset{6}{6}, \overset{7}{7}, \overset{8}{7}, \overset{9}{8}, \overset{10}{9}, \overset{11}{10}, \overset{12}{11}, \overset{13}{9}, \overset{14}{8}, \overset{15}{9}, \overset{16}{10} \}$

$$\underbrace{d^2(q_2, 4)}_{=2^2+1^2+1^2=6} \leq \epsilon^2/2 \quad \checkmark$$

$$\underbrace{d^2(q_2, 5)}_{=0^2+0^2+2^2=4} \leq \epsilon^2/2 \quad \checkmark$$

$$\underbrace{d^2(q_2, 6)}_{=1^2+1^2+2^2=6} \leq \epsilon^2/2 \quad \checkmark$$

.... remaining positions have sq. distance  $> \epsilon^2/2$

$S_2 = \{4, 5, 6\}$

...we also computed  $S_1 = \{1, 2, 3\}$

Positions in  $S_1$  correspond to candidate positions for  $q$ , positions in  $S_2$  correspond to candidate positions for  $q$ , **shifted by 3 time units on the right.**

→ therefore, positions for  $q$  that correspond to  $S_2$  are  $\{1, 2, 3\}$

$P = \text{merge}(S_1, S_2) = \{1, 2, 3\}$

We only need to check those positions for finding out subsequences for which  $q = \{3, 4, 3, 5, 6, 5\}$  is at most  $\epsilon=4$  distance away

$q_2 = \{5, 6, 5\}$

$\epsilon^2/2 = 8$

# Subsequence Matching: Example

$T = \{ \overset{1}{3}, \overset{2}{3}, \overset{3}{4}, \overset{4}{3}, \overset{5}{5}, \overset{6}{6}, \overset{7}{7}, \overset{8}{7}, \overset{9}{8}, \overset{10}{9}, \overset{11}{10}, \overset{12}{11}, \overset{13}{9}, \overset{14}{8}, \overset{15}{9}, \overset{16}{10} \}$

$$\underbrace{d^2(q, 1) = 0 + 1^2 + 1^2 + 2^2 + 1^2 + 1^2}_{\text{positions 2-7}} = 8 \leq \epsilon^2 \Rightarrow \text{result!}$$

$$\underbrace{d^2(q, 1) = 0 + 0 + 0 + 0 + 0 + 2^2}_{\text{positions 3-8}} = 4 \leq \epsilon^2 \Rightarrow \text{result!}$$

$$\underbrace{d^2(q, 1) = 0 + 0 + 0 + 0 + 0 + 2^2}_{\text{positions 4-9}} = 4 \leq \epsilon^2 \Rightarrow \text{result!}$$

$q = \{3, 4, 3, 5, 6, 5\}$

$\epsilon^2 = 16$


$P = \{1, 2, 3\}$

Access positions in P and verify results!

In fact, in this example, we could be sure that  $P = \{1, 2, 3\}$  is the result, without checking (WHY?)

# Subsequence Matching (cont'd)

- What happens if  $q$  is not a multiple of  $w$ ?



- Q: Can we use the same methodology using only  $q_1, q_2, q_3$ ?
  - Yes: We can prove that if  $D(q, s) \leq \epsilon$  then there for at least one of  $q_1, q_2, q_3$ ,  $D(q_i, s_i) \leq \epsilon/\sqrt{3}$ .
  - Based on the following truth:
    - If  $D(q_{1..n}, s_{1..n}) \leq \epsilon$  for  $n$ -length sequences  $q_{1..n}$  and  $s_{1..n}$ , then for any subsequence pair  $q_{i..j}, s_{i..j}$ ,  $1 \leq i \leq j \leq n$ , also  $D(q_{i..j}, s_{i..j}) \leq \epsilon$  holds

# Acknowledgement

---

- Some slides are taken/modified from

## **A Tutorial on Indexing and Mining Time Series Data**

**[http://www.cs.ucr.edu/%7Eeamonn/tutorial\\_on\\_time\\_series.ppt](http://www.cs.ucr.edu/%7Eeamonn/tutorial_on_time_series.ppt)**

- **Dr Eamonn Keogh**
- Computer Science & Engineering Department  
University of California - Riverside  
Riverside, CA 92521